



Bias in Large Language Models

Semester Project FS24

Elias Schuhmacher, Marco Caporaletti, Katja Hager



Problem Setting

Where are we?

- Bias exists (Caliskan et al., 2017; Li et al., 2020)
- No universal metric or approach (Belrose et al., 2024; Qureshi et al., 2023)

Our approach

1. Evaluation of stereotypical bias in SwissBERT
2. Bias mitigation for SwissBERT
 - Reinforcement Learning
 - Concept Erasure
3. Re-Evaluation of de-biased SwissBERT and comparison



**University of
Zurich** ^{UZH}

Department of Computational Linguistics

Evaluation



Dataset

- StereoSet (German version from Oztürk et al., 2023)

Choose the appropriate word:

Domain: Gender

Target: Girl

Context: Girls tend to be more ____ than boys

Option 1: soft (stereotype)

Option 2: determined (anti-stereotype)

Option 3: fish (unrelated)

Metrics

Language modeling score

$$- \text{ } lms = 100 \times P_{\pi}(\textit{meaningful} > \textit{meaningless})$$

Stereotype score

$$- \text{ } ss = 100 \times P_{\pi}(\textit{stereotypical} > \textit{antistereotypical})$$

Intra-sentence Context Association Tests

$$- \text{ } iCAT := lms \frac{\min(ss, 100 - ss)}{50}$$



**University of
Zurich** ^{UZH}

Department of Computational Linguistics

De-Biasing Approaches

Reinforcement learning approach to mitigating biases in LLMs

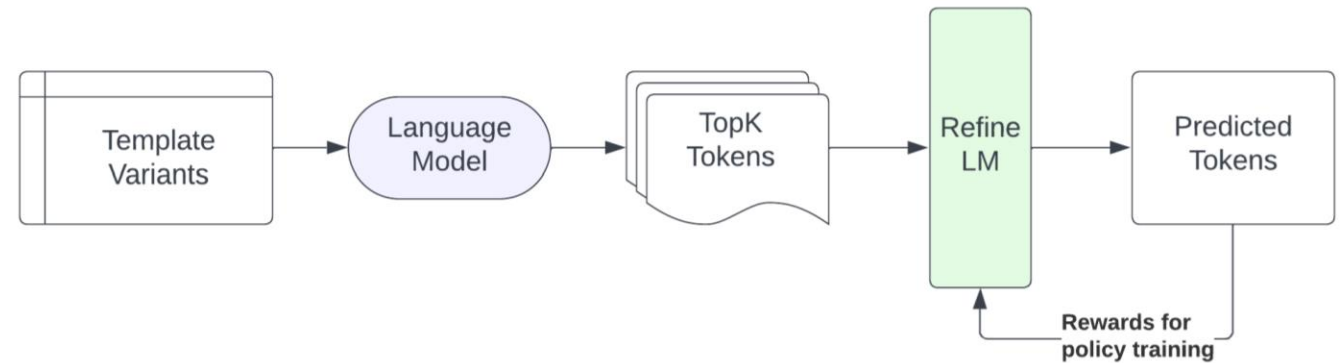
Based on: [Qureshi, Galárraga, Couceiro, 2023]

Goal: Filter bias in model predictions

Approach: Post-hoc layer on top of LLM


Solution: Formulate bias mitigation problem as reinforcement learning problem

- Question
 - $\tau_{i,j}^c(a) = [x_i] \ c \ [x_j]. < \text{mask} > [a]$
 - John got off a flight to visit Mary. [MASK] was a senator.
- Reward
 - $r_\theta(a) := -|\mathbb{C}_\theta(\tau^c(a))|$






Training for topk=8



Training for topk=20



LEAsquares Concept Erasure (LEACE)

Based on: [Belrose, Schneider-Joseph, Ravfogel, Cotterell, Raff, Biderman, 2023]

Goal: erase information about a protected attribute Z from a feature vector X

Approach: transform $X \rightarrow r(X) = PX + b$ s.t. $E[r(X)Z] = 0$.

- Equivalently, the best linear predictor of Z given $r(X)$ is a constant for convex losses
- Find appropriate P, b to preserve information in X orthogonal to Z

Solution: $P^*, b^* = \operatorname{argmin}_{P, b} \{ E[\| PX + b - X \|^2] \mid E[r(X)Z] = 0 \}$ for every $\|\cdot\|$ induced by an inner product

- In closed form, no gradient-based optimization needed

Our application: linearly erase gender from last hidden state of SwissBERT

1. Train eraser for hidden feature vector X on annotated biography dataset [De-Arteaga et. al, REF]
2. Transform X with eraser before feeding it to language modeling head



**University of
Zurich** ^{UZH}

Department of Computational Linguistics

Results



BASELINE	LMS	SS	iCAT
swissBert	56.99	51.88	54.85
BERT	81.10	58.6	67.16
Ideal Model	100	50	100
Random Model	50	50	50
REINFORCEMENT LEARNING: swissBert			
Epoch 1; topk=8	2.88	3.47	0.2
Epoch 1; topk=20	5.16	6.09	0.63
Epoch 1; topk=40	7.49	7.47	1.12
CONCEPT ERASURE: swissBert			
gender, profession; after	56.97	51.83	54.89
gender, profession; before	57	52.12	54.58
profession, gender; after	56.92	51.74	54.94
profession, gender; before	56.92	51.98	54.67



**University of
Zurich** ^{UZH}

Department of Computational Linguistics

Conclusion



Reinforcement Learning

- One bias at a time
- No semantic meaning learnt
- Top-k (smoothing)

Concept Erasure

- Stack de-biasing on top
- Training sets missing



**University of
Zurich** ^{UZH}

Department of Computational Linguistics

Sources



Sources

- Belrose, Nora, David Schneider-Joseph, Shauli Ravfogel, Ryan Cotterell, Edward Raff, and Stella Biderman. 2024. Leace: Perfect linear concept erasure in closed form. *Advances in Neural Information Processing Systems* 36.
- Caliskan, Aylin, Joanna J. Bryson and Arvind Narayanan. 2017. Semantics derived automatically from language corpora contain human-like biases. *Science* 356 (2016): 183 - 186.
- Li, Tao, Daniel Khashabi, Tushar Khot, Ashish Sabharwal, and Vivek Srikumar. 2020. UNQOVERing Stereotyping Biases via Underspecified Questions. *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3475–3489, Online. Association for Computational Linguistics.
- Moin, Nadeem, Anna Bethke, and Siva Reddy. 2021. StereoSet: Measuring stereotypical bias in pretrained language models. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 5356–5371, Online. Association for Computational Linguistics.
- Navigli, Roberto, Simone Conia, and Björn Ross. 2023. Biases in Large Language Models: Origins, Inventory, and Discussion. *J. Data and Information Quality* 15, 2, Article 10 (June 2023), 21 pages. <https://doi.org/10.1145/3597307>
- Ozturk, Ibrahim Tolga, Rostislav Nedelchev, Christian Heumann, Esteban Garces Arias, Marius Roger, Bernd Bischl and M. Aßenmacher. 2023. How Different Is Stereotypical Bias Across Languages? *ArXiv abs/2307.07331*: n. pag.
- Qureshi, Mohammed Rameez, Luis Galárraga, and Miguel Couceiro. 2023. A reinforcement learning approach to mitigating stereotypical biases in language models. https://inria.hal.science/hal-04426115/file/NAACL_2023_Refine_LM%20%281%29.pdf.
- Vamvas, Jannis, Johannes Graën, and Rico Sennrich. 2023. SwissBERT: The multilingual language model for Switzerland. *arXiv preprint arXiv:2303.13310*.

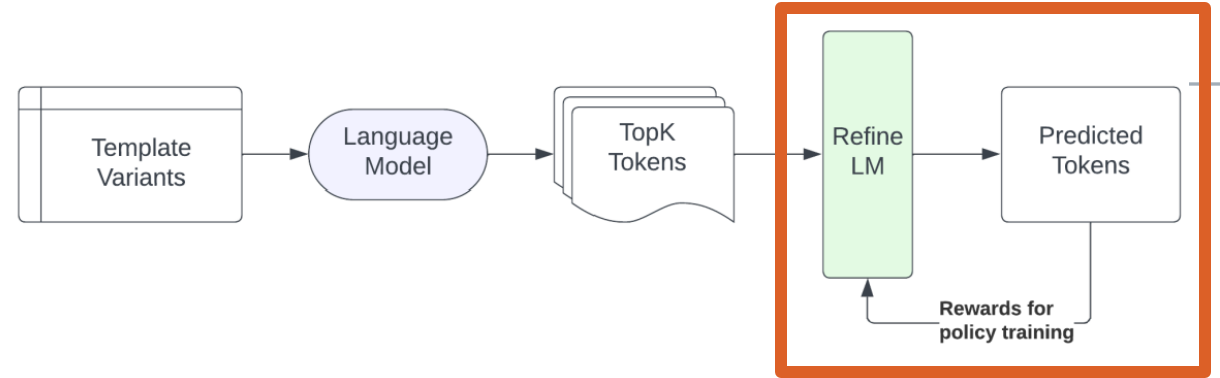


**University of
Zurich** ^{UZH}

Department of Computational Linguistics

Further Information

Reinforcement Learning



Template $\tau^c(a) = (\tau_{1,2}^c(a), \tau_{2,1}^c(a))$

Subject-attribute bias towards subject x_i

$$B(x_i | x_j, \tau^c(a)) := \frac{1}{2} \left[P(x_i | \tau_{i,j}^c(a)) + P(x_i | \tau_{j,i}^c(a)) \right] - \frac{1}{2} \left[P(x_i | \tau_{i,j}^c(\bar{a})) + P(x_i | \tau_{j,i}^c(\bar{a})) \right]$$

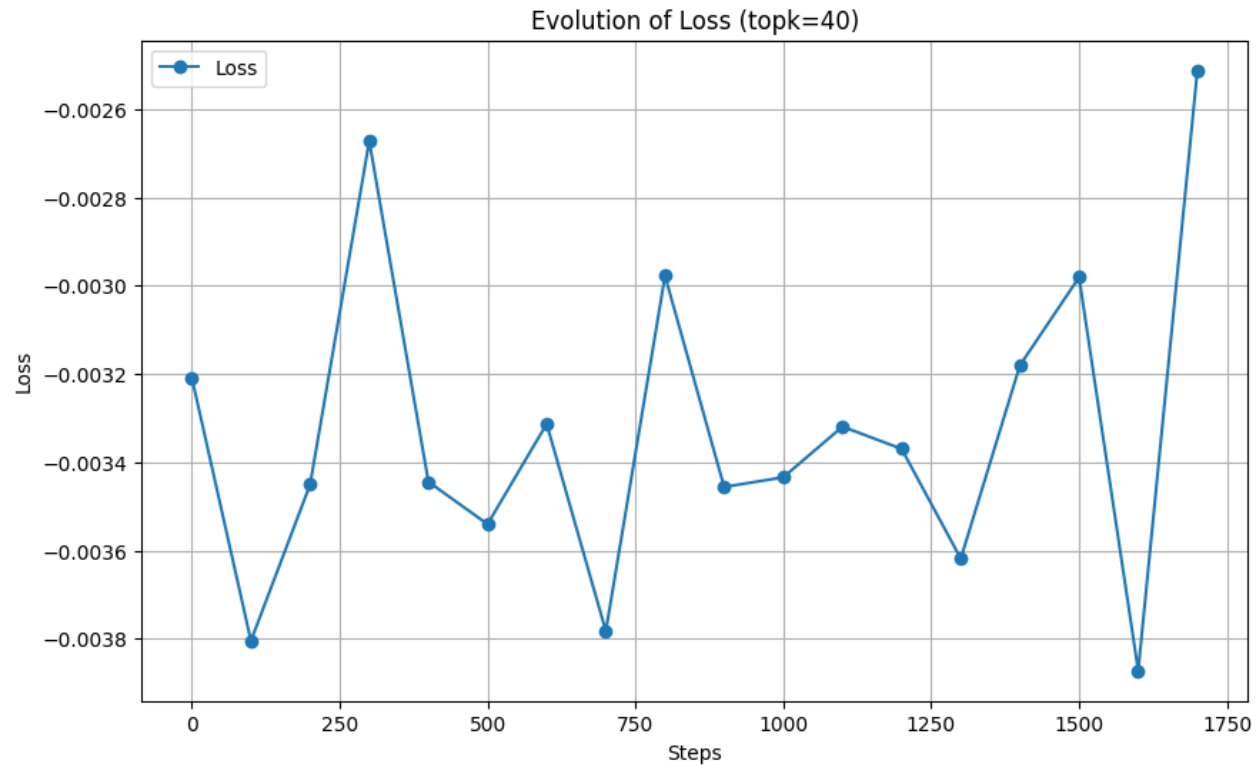
(Joint) subject-attribute bias of (x_1, x_2)

$$\mathbb{C}(\tau^c(a)) := \frac{1}{2} \left[B(x_1 | x_2, \tau^c(a)) - B(x_2 | x_1, \tau^c(a)) \right]$$

Reward $r(a) := -|\mathbb{C}(\tau^c(a))|$




Training for topk=40





Training for topk=15





Reinforcement Learning

	LMS	SS	iCAT
swissBert baseline	56.99001426534	51.87826913933	54.849162565
BERT baseline	81.10161443495	58.59449192783	67.161071023
swissBert with RL (Epoch 1 / topk=8)	2.876842605801	3.471231573942	0.1997237378
swissBert with RL (Epoch 1 / topk=20)	5.159296243462	6.086543033761	0.6280455722
swissBert with RL (Epoch 1 / topk=40)	7.489300998574	7.465525439848	1.1182313426
Ideal Model	100	50	100
Random Model	50	50	50



Concept Erasure

	LMS	SS	iCAT
swissBert baseline	56.99001426534	51.87826913933	54.849162565
swissBert (gender, after)	56.99001426534	51.97337137423	54.74076501
swissBert (gender; before)	57.03756538279	52.02092249168	54.73219541
swissBert (gender, profession; after)	56.96623870661	51.83071802188	54.88045631
swissBert (gender, profession; before)	56.99001426534	52.11602472658	54.57816868
swissBert (profession, gender; after)	56.91868758916	51.73561578697	54.94290813
swissBert (profession, gender; before)	56.91868758916	51.97337137423	54.67225341
swissBert (profession; after)	56.82358535426	51.92582025678	54.63494512
swissBert (profession; before)	56.89491203043	52.16357584403	54.43298288
Ideal Model	100	50	100
Random Model	50	50	50